

Author: Derek Riley

- Rule/prompt evaluation mechanism
 - Task-Specific Instructions
 - Arguments
 - Reusable Process Steps
 - Guided Examples and References
 - Explicit Output Requirements
 - Template-Based Naming
 - Error Handling and Edge Cases
 - Documentation
- Refactoring
 - Try adjusting the persona
 - *“As a senior software engineer who is an expert in maintainable java code...”*
 - Be specific with your goals
 - *Bad: “Refactor this”*
 - *Good: “Refactor the following function to improve its readability and maintainability (reduce repetition, use clearer variable names).”*
 - Refactor in steps- too much code can create distractions
- Debugging
 - Provide your intuition if you have it on the error
 - *“I suspect the root cause is due to the mishandling of XYZ edge case”*
 - Error messages can be copied/pasted, just ask the LLM that is the output you get and see how it responds
 - When something isn’t working, provide the input and what you are observing
 - Make sure the LLM has the context of what the code is supposed to be doing
 - LLMs can easily add debugging output
 - Performance/benchmarking output or progress is also an easy ask
 - Some errors don’t have error messages
 - Strategy: Ask the LLM to walk through the approach line-by-line
 - Example: *“Walk through this function line by line and track the value of total at each step. It’s not accumulating correctly – where does the logic go wrong?”*
 - Sometimes focusing the LLM on a block of code can lead it to find a bug
 - Consider changing LLM Roles
 - Ask the LLM to be a “code reviewer” to identify mistakes or bad practices
- System Testing
 - *“Act like a user of this application trying to break it. Create different inputs that test error checking. ”*
 - Best practice:
 - Have your system prompt instruct the Agent to write tests before writing code
 - Then only generate code that passes those tests
- Which LLM should I use?
 - [Artificialanalysis.ai](https://artificialanalysis.ai)

Anti-Patterns

- Vague prompting
 - Not enough context
 - Symptom: hallucination
 - Cure: context + clarity
- Overloaded prompt
 - Too much complexity, detail, scope
 - Symptom: incomplete, incorrect implementation
 - Cure: reduce scope/context, break down tasks
- Missing a call to action
 - Providing context without a clear call to action at the end
 - Symptom: Not doing what you expect
 - Cure: Make a clear, direct ask at the end of your prompt
- Buffet Coding, 1 of everything
 - Dumping all files into the context window, letting the LLM find the needle in the haystack
 - Symptom: Long generation times, lots of tokens consumed
 - Cure: Pre-select code chunks for inclusion in the context window
- Lacking directory structure
 - Letting the Agent create the directory with minimal or no guidance
 - Symptom: Files with strange names in strange locations
 - Cure: Create a folder, naming template, and any other structural guidance using standard naming conventions before prompting
- Vague completion expectations
 - Insufficient definition of “done”
 - Symptom: Incomplete/incorrect code
 - Cure: Make success criteria/tests very clear
- Ignoring explanations
 - Accepting edits without scanning provided justification
 - Symptom: Not reading output
 - Cure: Pay attention!
- Positional references
 - Referring to code as “above” or “previous”
 - Symptom: Incorrect/incomplete implementation
 - Cure: Use specific file/method/variable/etc names
- Over specificity
 - Too much detail, assuming AI doesn’t understand common concepts
 - Symptom: Overly detailed prompts with definitions of common concepts
 - Cure: Build familiarity with prompting to know what level of detail is necessary
- Fixing generated code
 - Generated code is buggy, close but not quite
 - Symptom: Your generated code is close but often requires manual tweaks
 - **Cure: Fix the process and context, not the code**